

**CS2106 Introduction to Operating Systems**  
**Lab 2 – Shell Scripting and Process Programming**  
**Release: Week of 9 September 2024**  
**Demo: Week of 16 September 2024**  
**Final Submission: Sunday 22 September 2024, 2359 hours**

Introduction

In this lab we will look at shells and at creating and managing processes. A shell is a program that runs on top of the operating system, and its main task is to interface the user with the operating system (OS). The familiar graphical environment in Windows and MacOS are both examples of graphical shells; they allow a user to interact with the OS using graphical elements like menu bars, context menus, text boxes, clickable links and images, etc.

Bash – Bourne Again Shell – is an example of a command-line shell, just like zsh in MacOS. Users interact with the OS by typing in commands. In this lab we will explore two aspects of programming command-line shells. In the first section we will look at how to do shell script programming on Bash, while in the second section we will look at how to create processes, run programs, do input and output redirection and pipes in C.

Instructions

Some points to note about this lab:

- a. This lab should be run using the slurm cluster, accessible via [xlog.comp.nus.edu.sg](http://xlog.comp.nus.edu.sg). Be sure to connect using the SoC VPN and NOT the NUS VPN.
- b. You may complete this lab on your own, or with ONE partner from any lab group. You do not need to partner with the same person with whom you completed Lab 1.
- c. There are two deliverables for this lab; a zip file that you can submit with your partner or individually IF you are doing the lab alone, and a practical demo that must be done individually.
- d. Only ONE copy of the submission zip file is to be submitted. Therefore, if you do the report with a partner, decide who should submit the report. DO NOT submit two copies.
- e. Use the enclosed AxxxxxY.docx answer book for your report, renaming it to the student ID of the submitter.
- f. Please indicate the student number, name, and group number of each person in the lab report.
- g. Create a ZIP file called AxxxxxY.zip with the following files (rename AxxxxxY to the student ID of the submitter):
  - The AxxxxxY.docx, appropriately renamed.
  - The process\_jobs.sh file from part b of Part 1.
  - The lab2p2f.c file from part b of Part 2.
- h. Upload your ZIP file to Canvas by SUNDAY, 22 September 2024, 2359 hours.**
- i. Some extra time is given for submission, but once the folder closes, NO SUBMISSIONS WILL BE ENTERTAINED AND YOU WILL RECEIVE 0 MARKS FOR THE LAB.**

## Part 1 – Bash Scripting

Bash scripting is an essential skill for anyone who works on servers running \*nix operating systems like Linux. It can automate many tasks. For example, you can write a Bash script to automatically compile your code, run unit tests if the compilation succeeds, and then push the code to a Github repository if the unit tests pass, while capturing the outputs of each stage to a file for later review.

Copy the Lab2Programs.zip file to xlog, and unzip it. Then switch to the “part1” directory.

This section introduces just the basics of shell scripting. For more details, please see <https://devhints.io/bash>

### a. Bash Script Basics

Before we start building a more interesting script, let’s go through some basics.

#### i. Creating a Hello World Script

Log on to xlog, and use your favorite editor and create a file called “hello.sh”. (the extension “sh” is conventional but unnecessary – you could have equally called it “hello.myhighfalutinshellscript” if you wanted. But please don’t.). Enter the following into “hello.sh”:

```
#!/bin/bash
echo "Hello world!" # Echo is similar to printf in C.
```

#### **Question 1.1 (1 mark)**

Ordinarily, comments in Bash scripts start with #. So, for example:

```
# This is a comment
ls -l          # This is also a comment
```

However, the first line of the Bash script is NOT a comment:

```
#!/bin/bash
```

(i) What is this line called? What is this line for?

(ii) What happens if this line is omitted or incorrectly specified?

Now exit your editor, and on the Bash command line, we convert it to an executable file by executing the following command:

```
chmod a+x ./hello.sh
```

This command sets the “executable” flag on hello.sh for “all” using the parameter “a+x”, thus everyone can execute hello.sh (Note: We will not use srin in this lab because it causes incorrect execution for some of the programs). To do so:

```
./hello.sh
```

You will see the output below:

```
pi@ubuntu-linux-20-04-desktop:~/cs2106/Labs/L02/soln/part1$ vim hello.sh
pi@ubuntu-linux-20-04-desktop:~/cs2106/Labs/L02/soln/part1$ chmod a+x ./hello.sh
pi@ubuntu-linux-20-04-desktop:~/cs2106/Labs/L02/soln/part1$ ./hello.sh
Hello world!
pi@ubuntu-linux-20-04-desktop:~/cs2106/Labs/L02/soln/part1$
```

**Note:** echo does not normally process ‘\n’ or other slash escape sequences. For example, if you did:

```
echo "\nHello world.\n"
```

You would get:

```
pi@ubuntu-linux-20-04-desktop:~/Desktop/labs/L02/soln/part1/autograder$ echo "\nHello world.\n"
\nHello world.\n
```

This is probably not what we want. To process escape sequences, specify the “-e” option when calling echo. E.g.

```
echo -e "\nHello world.\n"
```

You will now see that \n is properly processed:

```
pi@ubuntu-linux-20-04-desktop:~/Desktop/labs/L02/soln/part1/autograder$ echo -e "\nHello world.\n"
Hello world.
```

ii. Variables

Variables are very useful in Bash scripts, and can be used to store strings, numeric values, and even the outputs of programs. Use your favorite editor and create a file called "diff.sh", with the following lines to subtract 30 from 10, giving -20.

```
#!/bin/bash
x=10
y=30
z=$x-$y
echo "$x - $y = $z"
```

**NOTE:** In your assignment statements (e.g. `x=5`), it is VERY IMPORTANT that there is NO SPACE between the variable, the '=' and the value. Likewise in the line `z=$x-$y`, it is VERY IMPORTANT that THERE ARE NO SPACES in the statement. If you have spaces, you will get an error message like "x: command not found".

Use `chmod` to make this script executable, execute it and answer the following question:

**Question 1.2 (1 mark)**

The script produces the wrong result "`10 - 30 = 10-30`", instead of "`10 - 30 = -20`". Using Google or otherwise, fix the script so that it says "`10 - 30 = -20`".

- (i) Summarize in one line how you fixed the script.
- (ii) Explain why arithmetic operations in Bash require specific syntax.

Your currently hardcode the two inputs (10 and 30) in the script. Modify your script to take the two inputs from the command line. Copy and paste your script here.

Notice some things about the script above:

- (a) You assign to a variable using `=`, which is expected. However as mentioned earlier, there must not be any spaces in your assignment statement. For example, `x=5` is correct, but `x = 5` is wrong and will result in an error like "x: command not found"
- (b) Use `$<var name>` to access the value stored in `<var name>`. For example, we used `$x` and `$y` to access the values stored in `x` and `y`.
- (c) Notice that we can similarly access the values of the variables in the echo statement by using `$`.

Now let's look at how to capture the output of a program to a variable. On your xlog or other Linux shell session, type:

```
date +%A
```

This should print out the current day. For example:

```
pi@ubuntu-linux-20-04-desktop:~/cs2106/Labs/L02/soln/part1$ date +%A
Monday
```

To store this in a variable, we use the `$(.)` operator. Enter the following into your Bash shell (you do not need to write a script for this part):

```
day=$(date +%A)
echo "Today is %day."
```

You will see (this lab sheet was written on a Monday):

```
pi@ubuntu-linux-20-04-desktop:~/cs2106/Labs/L02/soln/part1$ day=$(date +%A)
pi@ubuntu-linux-20-04-desktop:~/cs2106/Labs/L02/soln/part1$ echo "Today is $day"
Today is Monday
```

### Question 1.3 (1 mark)

The “whoami” command returns the user ID of the current user:

```
[(base) mausamvora@Mausams-MacBook-Air Desktop % whoami
mausamvora
```

Now, using your favorite editor, create a shell script called “greet.sh” that greets the user, stating the user ID, the machine architecture, the kernel version, and the hostname in the following way:

```
(base) mausamvora@Mausams-MacBook-Air Desktop % ./greet.sh

Hello mausamvora, you are using a arm64 machine running kernel version
23.6.0 on Mausams-MacBook-Air.
```

**Hint:** You can type “man uname” (without the quotes) in Bash to learn more about the uname command, and how to extract the kernel version etc.

Copy and paste your code to your answer book.

### iii. Test Statements

Bash can test for certain conditions using the `[[.]]` operator. Some things you can do:

Test	Result
<code>[[ -z &lt;string&gt; ]]</code>	Tests if <string> is empty (i.e. "").
<code>[[ -n &lt;string&gt; ]]</code>	Tests if <string> is not empty.
<code>[[ &lt;string1&gt; == &lt;string2&gt; ]]</code>	Tests if <string1> is equal to <string 2>
<code>[[ &lt;string1&gt; != &lt;string 2&gt; ]]</code>	Tests if <string1> is not equal to <string 2>
<code>[[ num1 -eq num2 ]]</code>	(Numeric) Tests if num1 == num2
<code>[[ num1 -ne num2 ]]</code>	(Numeric) Tests if num1 != num2
<code>[[ num1 -lt num2 ]]</code>	(Numeric) Tests if num1 < num2
<code>[[ num1 -le num2 ]]</code>	(Numeric) Tests if num1 <= num2
<code>[[ num1 -gt num2 ]]</code>	(Numeric) Tests if num1 > num2
<code>[[ num1 -ge num2 ]]</code>	(Numeric) Tests if num1 >= num2
<code>[[ -e FILE ]]</code>	Tests if file exists
<code>[[ -d FILE ]]</code>	Tests if file is a directory
<code>[[ -s FILE ]]</code>	Tests if file size > 0
<code>[[ -x FILE ]]</code>	Tests if file is executable
<code>[[ FILE1 -nt FILE2 ]]</code>	Tests if FILE1 is new than FILE2
<code>[[ FILE1 -ot FILE2 ]]</code>	Tests if FILE1 is older than FILE2
<code>[[ FILE1 -ef FILE2 ]]</code>	Tests if FILE1 is the same as FILE2

Note the spaces after `[[` and before `]]`. They ARE important!

These tests can be use within `if..elif..else..fi` statements. Create a shell script called “comp.sh” and type in the following:

```
#!/bin/bash

echo "Enter the first number: "
read NUM1
echo "Enter the second number: "
read NUM2

if [[ NUM1 -eq NUM2 ]]; then
    echo "$NUM1 = $NUM2"
elif [[ NUM1 -gt NUM2 ]]; then
    echo "$NUM1 > $NUM2"
else
    echo "$NUM1 < $NUM2"
fi
```

Make `comp.sh` executable, and execute it. You can enter various numbers to play with it:

```

pi@ubuntu-linux-20-04-desktop:~/cs2106/Labs/L02/soln/part1$ ./comp.sh
Enter the first number:
2
Enter the second number:
3
2 < 3
pi@ubuntu-linux-20-04-desktop:~/cs2106/Labs/L02/soln/part1$ ./comp.sh
Enter the first number:
3
Enter the second number:
1
3 > 1
pi@ubuntu-linux-20-04-desktop:~/cs2106/Labs/L02/soln/part1$ ./comp.sh
Enter the first number:
5
Enter the second number:
5
5 = 5

```

Some things to note:

- You can read from the keyboard using “read”. The syntax is “read <varname>”, where <varname> is the variable that we want to store the read data to.
- The “if” statement has an odd syntax; you need a semi-colon after the [[..]]:

```

if [[ NUM1 -eq NUM2 ]]; then
    echo "$NUM1 = $NUM2"
else
    echo "$NUM1 != NUM2"
fi

```

#### iv. Loops

Bash supports both for-loops and while-loops. The for-loop is similar to Python’s. For example, to list all the files in the /etc directory, we could do:

```

for i in /etc/*; do
    echo $i
done

```

We get an output like this:

```

/etc/sos
/etc/speech-dispatcher
/etc/ssh
/etc/ssl
/etc/subgid
/etc/subgid-
/etc/subuid
/etc/subuid-
/etc/sudoers
/etc/sudoers.d
/etc/sysctl.conf

```

You can also iterate over a range:

```
for i in {1..5}; do
    echo $i
done
```

```
1
2
3
4
5
```

The while loop works pretty much the way you'd expect it to. For example:

```
i=0
while [[ $i -le 5 ]]; do
    echo $i
    let i=i+1
done;
```

We get:

```
0
1
2
3
4
5
```

You can read text files using the while loop! To print file.txt (provided in your zip file in the part1 directory), you can do:

```
cat file.txt | while read line; do
    echo $line
done
```

The “cat” command prints the contents of file.txt to the screen, but the “|” hijacks this output and sends it to the “read” command using a mechanism known as “pipe”. Recall that read line will read whatever is being piped in to the “line” variable.

```
This is a text file.
CS2106 Introduction to Operating Systems is a cool module!
```



v. Functions

You can declare a function called “func” this way. Parameters are accessed using \$1, \$2, etc. for the first parameter, second parameter, etc. Create a file called “func.sh” and type in the following:

```
#!/bin/bash
function func {
    echo "Called with $# parameters."
    sum=0
    for param in $@; do
        sum=$((sum + $param))
    done
    echo "The sum of all parameters is $sum"
    return $#
}

func 10 20 30
echo "The function returned: $?"
```

Make this file executable, then execute it. You will see:

```
Called with 3 parameters.
The sum of all parameters is 60
The function returned: 3
```

**Question 1.4 (1 mark)**

The following are special variables in Bash. What do they hold?

\$#, \$1, \$2, \$@, \$?

Can you change the above code, to give an output like this using special variables in Bash:

```
Called with 3 parameters.
Parameter 1 is hello
Parameter 2 is world
Parameter 3 is 83.6
139
```

vi. Miscellaneous Topics

Let’s look at a few additional features in Bash, which you may find useful.

a) Redirecting Output

You can redirect output to the screen (stdout) to a file. For example, if you wanted to capture the output of “ls” to a file:

```
ls > ls.out
```

Note that you can *append* to an output file by using >> instead of >. For example:

```
ls -l >> ls.out
```

Would append output to ls.out without overwriting previous content.

## b) Redirecting Input

You can also redirect input from the keyboard to a file. In the “part1” directory you have a file called “talk.c”. Compile and run it using:

```
gcc talk.c -o talk
./talk
```

This program echoes back whatever you type on the keyboard, prefacing it with “This is what I read:”

```
pi@ubuntu-linux-20-04-desktop:~/Desktop/labs/L02/soln/part1$ gcc talk.c -o talk
pi@ubuntu-linux-20-04-desktop:~/Desktop/labs/L02/soln/part1$ ./talk
hello
This is what I read: hello

this is a test
This is what I read: this is a test
```

There is also a file called file.txt. Enter the following command:

```
./talk < file.txt
```

As you can see in the output, the contents of file.txt are printed out by “talk” as though they were typed in on the keyboard:

```
pi@ubuntu-linux-20-04-desktop:~/Desktop/labs/L02/soln/part1$ ./talk < file.txt
This is what I read: This is a text file.

This is what I read: CS2106 Introduction to Operating Systems is a cool module!
```

### c) Getting the Result Returned by a Program

You have a file called **“slow.c”** in the **“part1”** directory. It reads a sentence from the command line and prints each word on a separate line. The source code is simple and shown below:

```
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <unistd.h>

int main(int ac, char **av) {
    if (ac != 2) {
        printf("\nPlease provide a sentence in quotes.\n");
        printf("Usage: %s \"<sentence>\"\n\n", av[0]);
        exit(-1);
    }

    char *sentence = strdup(av[1]);
    char *word = strtok(sentence, " ");

    int count = 0;
    while (word != NULL) {
        printf("%s\n", word);
        sleep(1);
        word = strtok(NULL, " ");
        count++;
    }

    printf("\nNumber of words: %d\n", count);

    // Try removing this exit to explore its effect
    exit(count);
}
```

Notice in particular the final statement **“exit(count);”**. It allows the program to return the total number of words in the sentence to the operating system. In this section we will see how to retrieve this value.

Compile and run it using the following instructions:

```
gcc slow.c -o slow
./slow "This is a test sentence"
```

You will see the following output:

```
(base) mausamvora@Mausams-MacBook-Air part1 % gcc slow.c -o slow
(base) mausamvora@Mausams-MacBook-Air part1 % ./slow "This is a test sentence"

This
is
a
test
sentence

Number of words: 5
```

Now enter the following command:

```
echo $?
```

#### Question 1.5 (1 mark)

Earlier we mentioned the “exit(count);” statement in slow.c.

(i) What do you see when you run “echo \$?”? How is echo \$? related to exit(count); in the program?

(ii) Modify the code:

- Remove the exit(count); statement and observe what happens when you run the program and check the value of echo \$?.

- Replace exit(count); with return count; and observe the results.

What do you notice about how these different approaches affect the return value and the behavior of the program?

#### d) Pipes

Assuming one program prints to the screen and another reads from the keyboard, you can channel the output of the first program to the input of the second using a mechanism called a “pipe”. Earlier we saw “slow.c” and “talk.c”. Enter the following command:

```
./slow "This is a test sentence" | ./talk
```

Your screen will appear to hang; wait for around 6-7 seconds and you will see:

```
(base) mausamvora@Mausams-MacBook-Air part1 % ./slow "This is a test sentence" | ./talk
This is what I read: This

This is what I read: is

This is what I read: a

This is what I read: test

This is what I read: sentence

This is what I read:

This is what I read: Number of words: 5
```

As you can see, the output of “slow” was channeled to the input of “talk”. Pipes are very useful mechanisms and we will see more of it later.

#### e) Running Programs Sequentially and In Parallel

Enter the following command:

```
./slow "This is a test sentence" ; ./slow "I have
increased the number of words in this sentence"
```

Now enter the following command:

```
./slow "This is a test sentence" & ./slow "I have
increased the number of words in this sentence"
```

#### Question 1.6 (1 mark)

(i) What happens when you run “./slow “This is a test sentence” ; ./slow “I have increased the number of words in this sentence””, and when you run “./slow “This is a test sentence” & ./slow “I have increased the number of words in this sentence””?

**What is the difference between “;” and “&”?**

(ii) You are given two tasks:

Task 1: Print 100 lines of a log file with a delay between each line.

Task 2: Send an email alert when the log contains a specific keyword.

Scenario: The system should print the log lines while scanning the log for the keyword at the same time. Once the keyword is detected, the system should immediately send the alert without waiting for all log lines to print. Based on the scenario, decide whether to use sequential or parallel execution for the tasks and justify your choice.

## b. Writing a Cool Script

We will now pull together everything you've learned in this section to write a script for batch processing multiple jobs. In **part1/batch**, the jobs are organized by students, each with their own directory inside the **jobs/** folder.

In the **part1/batch/jobs/** directory, you will find the following subdirectories:

Sub-Directory	Contents
A0183741Y/	Contains the student's program source files (.c, .h), and potentially some input files (.in)
A0281754H/	Contains the student's program source files (.c, .h), and potentially some input files (.in)
A0285757B/	Contains the student's program source files (.c, .h), and potentially some input files (.in)

Each folder contains the source code and input files for a job, which will be compiled, executed, and processed by your batch system.

A template shell script called "**process\_jobs.sh**" has already been created for you. This script is called with a single argument, which is the name of the program to be compiled and tested for each job. For example, if the name of the program is `sum`, you'd run the script like this:

```
./process_jobs.sh sum
```

Complete the template file. This is what your completed shell script should do:

- 1) If no argument is supplied to the script, it should output:

**Usage: ./process\_jobs.sh <filename>**

Similarly, if more than one argument is supplied to the script, it should output the same message.

- 2) If exactly one argument is supplied, the script should use `gcc` to compile the C source files inside each job folder. For example, for student A0183741Y, the script should compile the C files (`sum.c`, `utils.c`) like this:

```
gcc sum.c utils.c -o sum
```

The script should repeat this compilation process for each student folder in the **jobs/** directory.

- 3) For each input file (.in) found inside the student's directory, the compiled program should be executed, and the output should be written to an output file.
- 4) For example, if student A0183741Y has an input file `s1.in`, your script should run the program and redirect the input/output like this:

```
./sum < s1.in > s1.in.out
```

If there are multiple .in files (e.g., s1.in, s2.in), your script should use a loop to process each input file and generate the corresponding output file.

Your shell script SHOULD NOT hardcode the generation of the “.out” files, but must use a loop to iterate over every “.in” file to generate the corresponding “.out” file. While naming the output “f1.out” looks better than “f1.in.out”, it’s also harder so that’s not a requirement. ;)

- 5) Your script should iterate over every student directory in **jobs/**, and do:
- Compile the code using the appropriate .c files in the folder (see step 2 above).
  - Log an error in a **log.txt** file inside the student's folder if there is a compilation error.
  - Generate output files for each .in file found in the student folder (see step 3 and 4 above).
  - Measure the execution time (elapsed real time) for the program and log this information in **log.txt**.
  - Record whether the program encountered runtime errors (such as segmentation faults) and log these in **log.txt**.

- 6) Each “logs.txt” file should look somewhat like this:

```
Processing student job in jobs/A0183741Y
Compilation successful for jobs/A0183741Y
Execution time for jobs/A0183741Y/s1.in: 1 seconds
Execution completed without errors in jobs/A0183741Y
```

Specifically, the file should contain:

- A line showing the job is being processed
  - Compilation status (Success/Failure)
  - Execution time (elapsed real time) of the binary
  - Runtime errors encountered during execution (if any)
- 7) Generate a Summary Report:
- After processing all student folders, generate a **summary\_report.txt** file in the main **jobs/** directory that includes the following:
    - Total number of jobs processed.
    - Number of successful compilations.
    - Number of jobs that failed to compile.
    - Total runtime of all jobs.

- 8) Your “summary\_report.txt” file should look like this:

```
Summary Report
Total jobs processed: 3
Successful compilations: 2
Failed compilations: 1
Total runtime of all jobs: 1 seconds
```

**DEMO (4 marks)**

Run your shell script and show the output of **summary\_report.txt**, as well as **log.txt** and **output** files for all job folders to your TA. Points will be deducted for missing logs, incorrect statistics, or improperly processed jobs.

Some hints:

- 1) You should assume that **process\_jobs.sh** will be run in the directory that contains the **jobs/** folder.
- 2) Detect compilation errors by checking the return value of gcc.
- 3) Use the time command to measure execution time for each job. For how to get the elapsed real time, check “time --help”



## Part 2 – Playing with POSIX Process Calls

Change to the “part2” directory. In this section we will learn how to create processes, how to parallelize processes, and how to redirect input and output, and how to set up pipes between processes.

### a. Introduction to Process Management in C

#### i. Creating a new Process

You can spawn a new process by using “fork”. Open the “lab2p2a.c” file with your favorite editor.

Examine the code:

```
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
#include <sys/wait.h>
#include <string.h>

int slow(char *name, int time, char *sentence) {
    char *word = strtok(sentence, " ");
    int count = 0;
    while(word != NULL) {
        printf("%s: %s\n", name, word);
        sleep(time); // Simulate a delay for each word
        word = strtok(NULL, " ");
        count++;
    }
    printf("\n%s: Number of words: %d\n", name, count);
    return count;
}

int main() {
    int id;
    char sentence1[] = "This is a test sentence";
    char sentence2[] = "I have increased the number of words in this sentence";

    if((id = fork()) != 0) {
        int stat;
        int my_id = getpid();
        int parent_id = getppid();
        printf("\nI am the parent.\n");
        printf("My ID is %d\n", my_id);
        printf("My parent's ID is %d\n", parent_id);
        printf("My child's ID is %d\n", id);
        slow("Parent", 1, sentence1);
        printf("\nWaiting for child to exit.\n");
        wait(&stat);
        printf("CHILD HAS EXITED WITH STATUS %d\n", WEXITSTATUS(stat));
    } else {
        id = getpid();
        int parent_id = getppid();
        printf("\nI am the child.\n");
        printf("My ID is %d\n", id);
        printf("My parent's ID is %d\n\n", parent_id);
        int word_count = slow("Child", 2, sentence2);
        exit(word_count);
    }
}
```

Some explanation of this code:

- We `#include <stdlib.h>` to bring in “exit”, which the child will use to return a value.
- We `#include <unistd.h>` to bring in the prototype for “fork”, used to spawn a new process, and `#include <sys/wait.h>` to bring in “wait”.
- The **slow** function splits the sentence into words and prints each word with a delay, simulating a "slow" process, similar to the `slow.c` program we saw earlier. The parent process prints a number every second and the child process prints a number every two seconds.
- The “fork” function call creates a new process:
  - o “fork” returns 0 to the child process, and the child’s process ID (PID) – a non-zero value – to the parent.
  - o A process can get its own PID using `getpid()`, and its parent’s PID using `getppid()`.
- The child returns a value of “word\_count” using “exit”.
- The parent calls “wait” to wait for the child to finish, stores the child’s return value into “stat”, then calls the `WEXITSTATUS` macro to extract the returned value from “stat”. You can find out more about the various exit status macros here: [https://www.gnu.org/software/libc/manual/html\\_node/Process-Completion-Status.html](https://www.gnu.org/software/libc/manual/html_node/Process-Completion-Status.html)

Compile and execute the above code using:

```
gcc lab2p2a.c -o lab2p2a
./lab2p2a
```

Observe the results and answer the question below:

**Question 2.1 (1 mark)**

Are the parent and child processes executing in concurrently? How do you know this?

Modify the code:

- Try running the program with and without wait() and describe how removing wait() affects the execution order and the interleaving of output between the parent and child processes.

**Question 2.2 (1 mark)**

You can see that that the parent process also has a parent. Who is the parent's parent? **Hint:** The "ps" command shows the processes running in the current shell together with their PIDs. Execute the following two commands on two different terminals.

First Terminal: ./lab2p2a

Second Terminal: ps | grep "lab2p2a"

Discuss briefly the process hierarchy in operating systems.

ii. Accessing Arguments and Environment Variables in C

Use your favorite editor to open "lab2p2b.c". Examine the code:

```
#include <stdio.h>

int main(int ac, char **av, char **vp) {
    printf("ac = %d\n", ac);
    printf("Arguments:\n");

    int i;

    for(i=0; i<ac; i++)
        printf("Arg %d is %s\n", i, av[i]);

    i=0;
    while(vp[i] != NULL) {
        printf("Env %d is %s\n", i, vp[i]);
        i++;
    }
}
```

Compile and execute the program using the following commands, observe the outputs and answer the questions that follow:

```
gcc lab2p2b.c -o lab2p2b
```

```
./lab2p2b
./lab2p2b hello world
./lab2p2b this is a test
```

### Question 2.3 (1 mark)

What do “ac”, “av” and “vp” contain?

### iii. Loading and Executing A Program

In C we can load and execute a program using the “exec\*” family of system calls. The table below shows the different versions. You must #include <unistd.h> to use these functions.

Version	What it Does
<code>execl(const char *path, const char *arg1, char *arg2, const char arg3, ...)</code>	<p>Executes command shown in “path”. Arguments to command are listed individually, terminated with a NULL.</p> <p>E.g.</p> <pre>execl("/bin/ls", "ls", "-l", NULL);</pre> <p>Note you must specify the full path to “ls”. Also conventionally the first argument is always the name of the program you are running.</p>
<code>execlp(const char *file, const char *arg1, char *arg2, const char arg3, ...)</code>	<p>Like execl, except that if you do not specify the full path to the command to run, the OS will search for the command in all directories specified in the PATH environment variable.</p> <p>E.g.</p> <pre>execlp("ls", "ls", "-l", NULL);</pre>
<code>execv(const char *path, char *const argv[]);</code>	<p>Like execl, except that the arguments are specified in an array instead of individually. The last element of the array must be NULL:</p> <p>E.g.</p> <pre>char *args[] = {"ls", "-l", NULL}; execv("/bin/ls", args);</pre>
<code>execvp(const char *file, char *const argv[])</code>	<p>Like execv, except that if you do not specify the full path to the command to run, the OS will search for the command in all directories specified in the PATH environment variable.</p>

	<p>E.g.</p> <pre>char *args[] = {"ls", "-l", NULL}; execvp("ls", args);</pre>
--	---

There are also `execle` and `execve` function calls which pass in environment variables, and we will ignore these here.

Since all the `exec*` functions replace the current process image with the process image of the program being run, and is conventionally run within a `fork()`. Again use your favorite editor to open the “lab2p2c.c” file, and examine the code:

```
#include <unistd.h>
#include <stdio.h>
#include <sys/wait.h>

int main() {
    if(fork() == 0) {
        execlp("cat", "cat", "file.txt", NULL);
    }
    else
        wait(NULL);
}
```

#### Question 2.4 (1 mark)

Change the code to use `execvp` instead of `execlp`. Cut and paste your new code here and explain what you’ve done.

#### iv. Redirecting Input and Output

In the previous section we saw how we can redirect input and output using “<” and “>” respectively. Now we will see how to do this programmatically:

- (a) You will see a “talk.c” program in the “part2” directory. This is exactly the same as the “talk.c” program in part 1. Compile it, naming the output executable “talk”.

(b) Open lab2p2d.c, and examine the code:

```
#include <stdio.h>
#include <unistd.h>
#include <fcntl.h>
#include <stdlib.h>
#include <sys/wait.h>

int main() {
    int fp_in = open("./file.txt", O_RDONLY);
    int fp_out = open("./talk.out", O_CREAT | O_WRONLY);

    if(fork() == 0) {
        dup2(fp_in, STDIN_FILENO);
        dup2(fp_out, STDOUT_FILENO);
        execlp("./talk", "talk", (char *) 0);
        close(fp_in);
        close(fp_out);
    }
    else
        wait(NULL);
}
```

A few things to note:

- We are going to use library calls to execute the equivalent of:

```
./talk < file.txt > talk.out
```

- We are using the more primitive “open” and “close” operations to operations to open file.txt for reading (O\_RDONLY). We also create a new file called talk.out for writing (O\_CREAT | O\_WRONLY). We use these instead of fopen and fclose and open and close will create file descriptors in the form that we need here.

Compile and run the program using:

```
gcc lab2p2d.c -o lab2p2d
./lab2p2d
cat talk.out
```

(Note: If you get “cat: talk.out: permission denied”, chmod the permissions for talk.out to 0400)

Notice that our program has done exactly ./talk < file.txt > talk.out as mentioned earlier.

**Question 2.5 (1 mark)**

- (i) What does “dup2” do? Why do we use “dup2” here?
- (ii) The program calls close() after execlp(). What happens to these file descriptors after the program calls execlp and how does the close() call affect them?

**v. Pipes**

A “pipe” is a byte-oriented communication mechanism between two processes using two file handles; the first is for reading, and the second is for writing.

Open the lab2p2e.c file and you will see the following code:

```
#include <stdio.h>
#include <unistd.h>
#include <fcntl.h>
#include <string.h>
#include <sys/wait.h>

int main() {
    int p[2];
    char str[] = "Hello this is the parent.";

    // This creates a pipe. p[0] is the reading end,
    // p[1] is the writing end.

    if(pipe(p) < 0)
        perror("lab2p2e: ");

    // We will send a message from father to child
    if(fork() != 0) {
        close(p[0]); // The the end we are not using.
        write(p[1], str, strlen(str));
        close(p[1]);
        wait(NULL);
    }
    else
    {
        char buffer[128];

        close(p[1]); // Close the writing end
        read(p[0], buffer, 127);
        printf("Child got the message \"%s\"\n", buffer);
        close(p[0]);
    }
}
```

We note some points of this code:

- Most POSIX calls will return a value of less than 0 if there's been an error. In this code we check for errors for the first time in these lines:

```
if(pipe(p) < 0)
    perror("lab2p2e: ");
```

The “perror” call prints the cause of the error to the screen, preceding it with the string “lab2p2e:”.

- The “pipe” call takes in an integer array p of two elements. The p[0] element is the reading end of the pipe, and the p[1] element is the writing end of the pipe.
- Just as we used “open” and “close” instead of “fopen” and “fclose” in the previous example, here we again use the more primitive “read” and “write” operations. You can see from the code how these are used.
- We always close the end we are not using. For example, the parent is not using the reading end and thus closes it, and likewise with the child.

Run the code by doing:

```
gcc lab2p2e.c -o lab2p2e
./lab2p2e
```

Observe that the child has received the parent's message.

#### Question 2.6 (1 mark)

What will happen if we do not close the unused pipe ends?

#### b. Piping Between Commands

We will now attempt something more challenging; we will write a program that pipes the output of one program to the input of another.

- (a) Compile the talk.c program if it's not already compiled. Name the executable file “talk”.
- (b) Compile the slow.c program, also in the part2 directory. Name the executable file “slow”.
- (c) Open the file “lab2p2f.c”, and write the code that does the equivalent of the following using C pipes, redirection, fork and exec\* as you've learnt before.

```
./slow "This is a test sentence" | ./talk > results.out
```



Some hints:

- ./slow will take about 7 seconds to run, so your computer will appear frozen for that time even when it is correctly done.
- You can execute `"./slow "This is a test sentence" | ./talk"`, then do:

```
gcc lab2p2f.c -o lab2p2f
./lab2p2f
cat results.out
```

You will see an output similar what you saw above.

- The essential thing is that you need to run ./slow and ./talk as two processes, then somehow redirect the output of one program to the writing end of the pipe, and the input of the other program to the reading end of the pipe.
- Don't forget to redirect the final result to "results.out".

**Question 2.7 (1 mark)**

Explain briefly how you set up the pipe between ./slow and ./talk